

# Learning Notes

DAVID

Reinforcement Learning Notes

## 1 Part1

### 1.1 Q3

**Markov Decision Process.** Markov Decision Process (MDP)  $\mathcal{M}_{\text{Env}} := (\mathcal{S}, \mathcal{A}, P_0, P, R)$  with states  $s \in \mathcal{S}$ , actions  $a \in \mathcal{A}$ , initial state distribution  $P_0$ , transition probabilities  $P$ , and reward  $R$ .

At each timestep  $t$ , the agent (e.g., robot) observes the state  $s_t \in \mathcal{S}$ , takes an action  $a_t \sim \pi(a_t | s_t) \in \mathcal{A}$ , transitions to the next state  $s_{t+1}$  according to  $s_{t+1} \sim P(s_{t+1} | s_t, a_t)$  while receiving the reward  $R(s_t, a_t)$ <sup>2</sup>. Fixing the MDP  $\mathcal{M}_{\text{Env}}$ , we let  $\mathbb{E}^\pi$  (resp.  $\mathbb{P}^\pi$ ) denote the expectation (resp. probability distribution) over trajectories,  $\tau = (s_0, a_0, \dots, s_T, a_T)$  with length  $T + 1$ , with initial state distribution  $s_0 \sim P_0$  and transition operator  $P$ . We aim to train a policy to optimize the cumulative reward, discounted by a function  $\gamma(\cdot)$ , such that the agent receives:

$$\mathcal{J}(\pi_\theta) = \mathbb{E}^{\pi_\theta, P_0} \left[ \sum_{t \geq 0} \gamma(t) R(s_t, a_t) \right].$$

**Policy optimization.** The policy gradient method (e.g., REINFORCE [85]) allows for improving policy performance by approximating the gradient of this objective w.r.t. the policy parameters:

$$\nabla_\theta \mathcal{J}(\pi_\theta) = \mathbb{E}^{\pi_\theta, P_0} \left[ \sum_{t \geq 0} \nabla_\theta \log \pi_\theta(a_t | s_t) r_t(s_t, a_t) \right], \quad r_t(s_t, a_t) := \sum_{\tau \geq t} \gamma(\tau) R(s_\tau, a_\tau),$$

where  $r_t$  is the discounted cumulative future reward from time  $t$  (more generally,  $r_t$  can be replaced by a Q-function estimator [76]),  $\gamma$  is the discount factor that depends on the time-step, and  $\nabla_\theta \log \pi_\theta(a_t | s_t)$  denotes the gradient of the logarithm of the likelihood of  $a_t | s_t$ . To reduce the variance of the gradient estimation, a state-value function  $\hat{V}^{\pi_\theta}(s_t)$  can be learned to approximate  $\mathbb{E}[r_t]$ . The estimated advantage function  $\hat{A}^{\pi_\theta}(s_t, a_t) := r_t(s_t, a_t) - \hat{V}^{\pi_\theta}(s_t)$  substitutes  $r_t(s_t, a_t)$ .

**Diffusion models.** A denoising diffusion probabilistic model (DDPM) [49, 29, 71] represents a continuousvalued data distribution  $p(\cdot) = p(x^0)$  as the reverse denoising process of a forward noising process  $q(x^k | x^{k-1})$  that iteratively adds Gaussian noise to the data. The reverse process is parameterized by a neural network

<sup>1</sup> More generally, we can view our environment as a Partially Observed Markov Decision Process (POMDP) where the agent's actions depend on observations  $o$  of the states  $s$  (e.g., action from pixels). Our implementation applies in this setting, but we omit additional observations from the formalism to avoid notional clutter.

---

Author's Contact Information: DAVID.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

*Conference acronym 'XX, June 03–05, 2018, Woodstock, NY*

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/18/06

<https://doi.org/XXXXXXXX.XXXXXXXX>

Table 1. Notations Table.

Notations	Descriptions
$t$	timestamps
$\mathcal{M}_{\text{Env}}$	MDP
$s \in \mathcal{S}$	states
$a \in \mathcal{A}$	actions
$P_0$	initial state distribution
$P$	transition probabilities
$R$	reward

<sup>2</sup> For simplicity, we overload  $R(\cdot, \cdot)$  to denote both the random variable reward and its distribution.  $\varepsilon_\theta(x_k, k)$ , predicting the added noise  $\varepsilon$  that converts  $x_0$  to  $x_k$  [29]. Sampling starts with a random sample  $x^K \sim \mathcal{N}(0, \mathbf{I})$  and iteratively generates the denoised sample:

$$x^{k-1} \sim p_\theta(x^{k-1} | x^k) := \mathcal{N}(x^{k-1}; \mu_k(x^k, \varepsilon_\theta(x^k, k)), \sigma_k^2 \mathbf{I}).$$

Above,  $\mu_k(\cdot)$  is a fixed function, independent of  $\theta$ , that maps  $x^k$  and predicted  $\varepsilon_\theta$  to the next mean, and  $\sigma_k^2$  is a variance term that abides by a fixed schedule from  $k = 1, \dots, K$ . We refer the reader to Chan [10] for an in-depth survey.

**Diffusion models as policies.** Diffusion Policy (DP; see [15]) is a policy  $\pi_\theta$  parameterized by a DDPM which takes in  $s$  as a conditioning argument, and parameterizes  $p_\theta(a^{k-1} | a^k, s)$  as in (3.3). DPs can be trained via behavior cloning by fitting the conditional noise prediction  $\varepsilon_\theta(a^k, s, k)$  to predict added noise. Notice that unlike more standard policy parameterizations such as unimodal Gaussian policies, DPs do not maintain an explicit likelihood of  $p_\theta(a^0 | s)$ . In this work, we adopt the common practice of training DPs to predict an action chunk - a sequence of actions a few time steps (denoted  $T_a$ ) into the future - to promote temporal consistency. For fair comparison, our non-diffusion baselines use the same chunk size.

As observed in [16, 6] and [59], a denoising process can be represented as a multi-step MDP in which policy likelihoods can be obtained directly. We extend this formalism by embedding the Diffusion MDP into the environmental MDP, obtaining a larger "Diffusion Policy MDP" denoted  $\mathcal{M}_{\text{DP}}$ , visualized in Fig. 3. Below, we use the notation  $\delta$  to denote a Dirac distribution and  $\otimes$  to denote a product distribution.

Recall the environment MDP  $\mathcal{M}_{\text{Env}} := (\mathcal{S}, \mathcal{A}, P_0, P, R)$  in Section 3. The Diffusion MDP  $\mathcal{M}_{\text{DP}}$  uses indices  $\bar{t}(t, k) = tK + (K - k - 1)$  corresponding to  $(t, k)$ , which increases in  $t$  but (to keep the indexing conventions of diffusion) decreases lexicographically with  $K - 1 \geq k \geq 0$ . We write states, actions and rewards as,

$$\bar{s}_{\bar{t}(t,k)} = (s_t, a_t^{k+1}), \quad \bar{a}_{\bar{t}(t,k)} = a_t^k, \quad \bar{R}_{\bar{t}(t,k)}(\bar{s}_{\bar{t}(t,k)}, \bar{a}_{\bar{t}(t,k)}) = \begin{cases} 0 & k > 0 \\ R(s_t, a_t^0) & k = 0 \end{cases},$$

where the bar-action at  $\bar{t}(t, k)$  is the action  $a_t^k$  after one denoising step. Reward is only given at times corresponding to when  $a_t^0$  is taken. The initial state distribution is  $\bar{P}^0 = P_0 \otimes \mathcal{N}(0, \mathbf{I})$ , corresponding to  $s_0 \sim P_0$  is the initial distribution from the environmental MDP and  $a_0^K \sim \mathcal{N}(0, \mathbf{I})$  independently. Finally, the transitions are

$$\bar{P}(\bar{s}_{\bar{t}+1} | \bar{s}_{\bar{t}}, \bar{a}_{\bar{t}}) = \begin{cases} (s_t, a_t^k) \sim \delta_{(s_t, a_t^k)} & \bar{t} = \bar{t}(t, k), k > 0 \\ (s_{t+1}, a_{t+1}^K) \sim P(s_{t+1} | s_t, a_t^0) \otimes \mathcal{N}(0, \mathbf{I}) & \bar{t} = \bar{t}(t, k), k = 0 \end{cases}.$$

That is, the transition moves the denoised action  $a_t^k$  at step  $\bar{t}(t, k)$  into the next state when  $k > 0$ , or otherwise progresses the environment MDP dynamics with  $k = 0$ . The pure noise  $a_t^K$  is considered part of the environment when transitioning at  $k = 0$ . In light of (3.3), the policy in  $\mathcal{M}_{\text{DP}}$  takes the form

$$\bar{\pi}_\theta(\bar{a}_{\bar{t}(t,k)} \mid \bar{s}_{\bar{t}(t,k)}) = \pi_\theta(a_t^k \mid a_t^{k+1}, s_t) = \mathcal{N}\left(a_t^k; \mu\left(a_t^{k+1}, \varepsilon_\theta\left(a_t^{k+1}, k+1, s_t\right)\right), \sigma_{k+1}^2 \mathbf{I}\right).$$

Fortunately, (4.1) is a Gaussian likelihood, which can be evaluated analytically and is amenable to the policy gradient updates (see also [59] for an alternative derivation):

$$\nabla_\theta \bar{\mathcal{J}}(\bar{\pi}_\theta) = \mathbb{E}^{\bar{\pi}_\theta, \bar{P}, \bar{P}^0} \left[ \sum_{\bar{t} \geq 0} \nabla_\theta \log \bar{\pi}_\theta(\bar{a}_{\bar{t}} \mid \bar{s}_{\bar{t}}) \bar{r}(\bar{s}_{\bar{t}}, \bar{a}_{\bar{t}}) \right], \quad \bar{r}(\bar{s}_{\bar{t}}, \bar{a}_{\bar{t}}) := \sum_{\tau \geq \bar{t}} \gamma(\tau) \bar{R}(\bar{s}_\tau, \bar{a}_\tau).$$

Evaluating the above involves sampling through the denoising process, which is the usual "forward pass" that samples actions in Diffusion Policy; as noted above, the initial state can be sampled from the environment via  $\bar{P}^0 = P_0 \otimes \mathcal{N}(0, \mathbf{I})$ , where  $P_0$  is from the environment MDP. 4.2 Instantiating DPPO with Proximal Policy Optimization

We apply Proximal Policy Optimization (PPO) [70, 18, 32, 1], a popular improvement of the vanilla policy gradient update.

**Definition 4.1 (Generalized PPO).** Consider a general MDP. Given an advantage estimator  $\hat{A}(s, a)$ , the PPO update is given by [70] the sample approximation to where  $\varepsilon$ , the clipping ratio, controls the maximum magnitude of policy change from the previous policy. We instantiate PPO in our diffusion MDP with  $(s, a, t) \leftarrow (\bar{s}, \bar{a}, \bar{t})$ . Our advantage estimator takes a specific form that respects the two-level nature of the MDP: let  $\gamma_{\text{ENV}} \in (0, 1)$  be the environment discount and  $\gamma_{\text{DENOISE}} \in (0, 1)$  be the denoising discount. Consider the environment-discounted return:

$$\bar{r}(\bar{s}_{\bar{t}}, \bar{a}_{\bar{t}}) := \sum_{t' \geq \bar{t}} \gamma_{\text{ENV}}^{t'} \bar{r}(\bar{s}_{\bar{t}(t', 0)}, \bar{a}_{\bar{t}(t', 0)}), \quad \bar{t} = \bar{t}(t, k),$$

## 1.2 Q4

## 1.3 Q6

## References